

Active Inference

Practical Part: Computational Phenotyping in Psychiatry

Philipp Schwartenbeck

Centre for Cognitive Neuroscience, Salzburg

Wellcome Trust Centre for Neuroimaging, London

Computational Psychiatry Course 2016, TNU, Zurich

Active Inference Essentials

Belief-based (choice) behaviour

- Agents build **generative models** of their environment.
 - I.e. joint probability distributions over observations y and causes θ : $P(y, \theta) = P(y|\theta) \cdot P(\theta)$
- Agents invert these models to infer latent variables, including current states, policies and the precision of beliefs

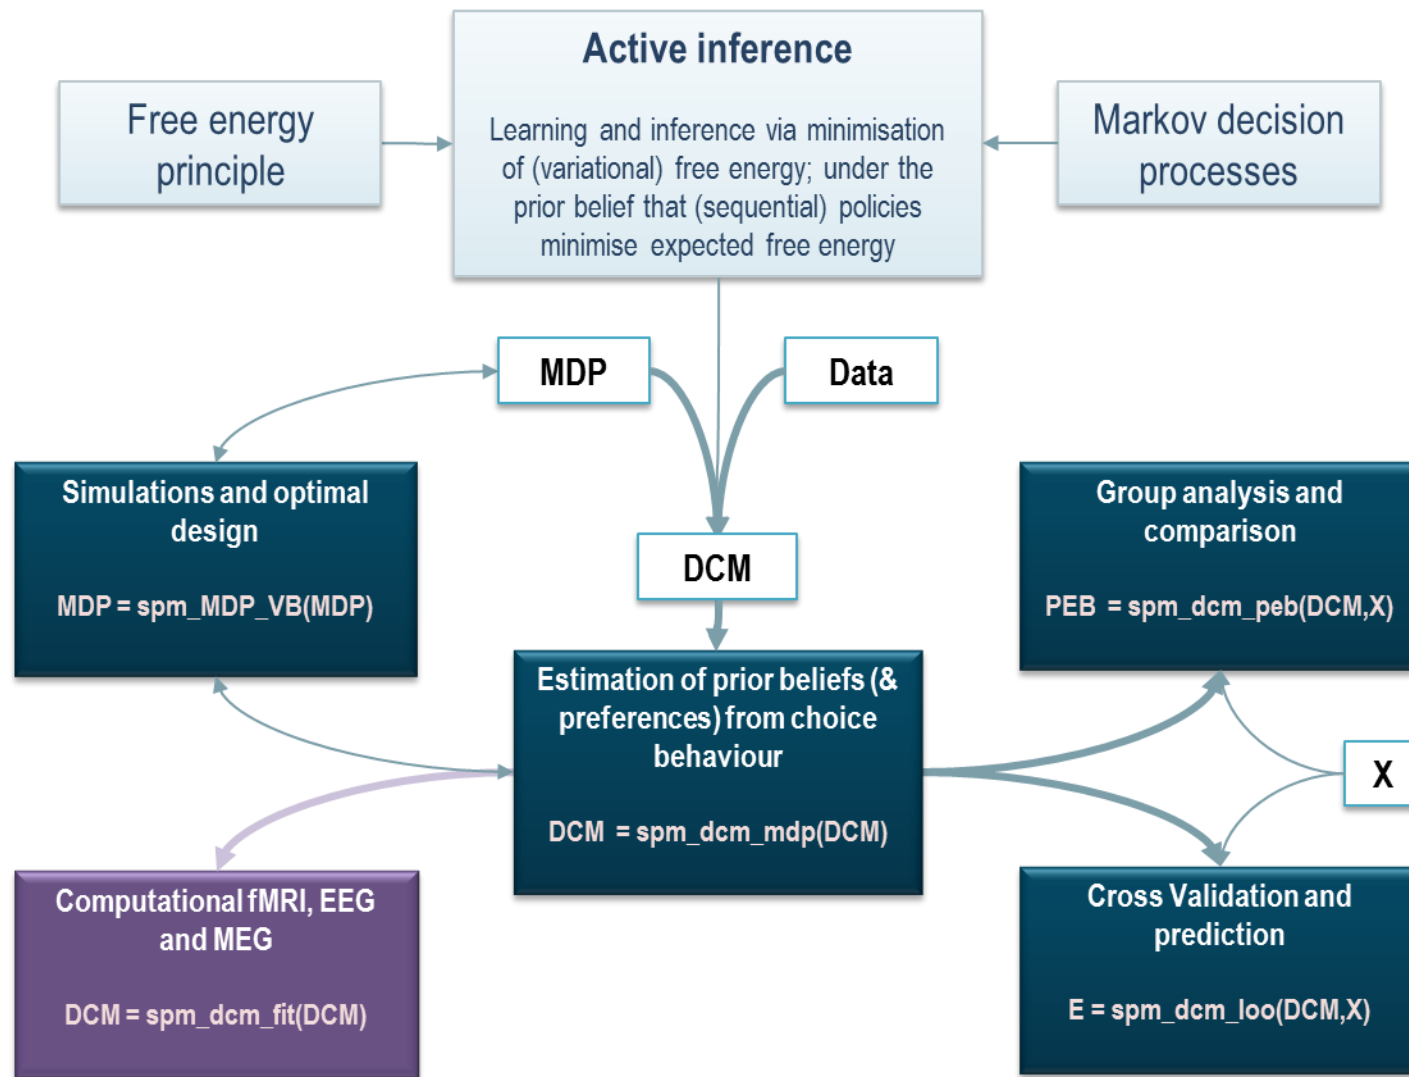
Preferences are cast as beliefs (expectations) about future states

- Agents try to fulfil their expectations
- Rewarding outcomes = expected outcomes

Objective function (cf., *computational level*, Marr): **minimise surprise**


- That means fulfilling expectations and maximising model evidence

Computational Phenotyping: Overview



Computational Phenotyping: Overview


- I. ABC of choice behaviour: model a task that involves exploratory and exploitative behaviour
- II. Simulate data and invert model to obtain subject-specific parameters
- III. Simulate group effect and recover group-specific parameters
- IV. Perform hierarchical Bayesian inference on group effects



Computational Phenotyping: Task

Two-step maze task

- A rat needs to obtain a reward in the left or right arm of a T-shaped maze
- It starts in the middle and can decide to go either left or right – or sample a cue at the bottom first
- The cue tells the rat with 95% validity where the reward is located
- The left and right arm are *absorbing states* (the rat cannot sample both)




Thus, the rat needs to solve the exploitation-exploration dilemma

- In case of high uncertainty, it should sample the cue first

Subjective Generative Model


Control states

$u \in U$





Hidden states

$s \in S$



Outcomes

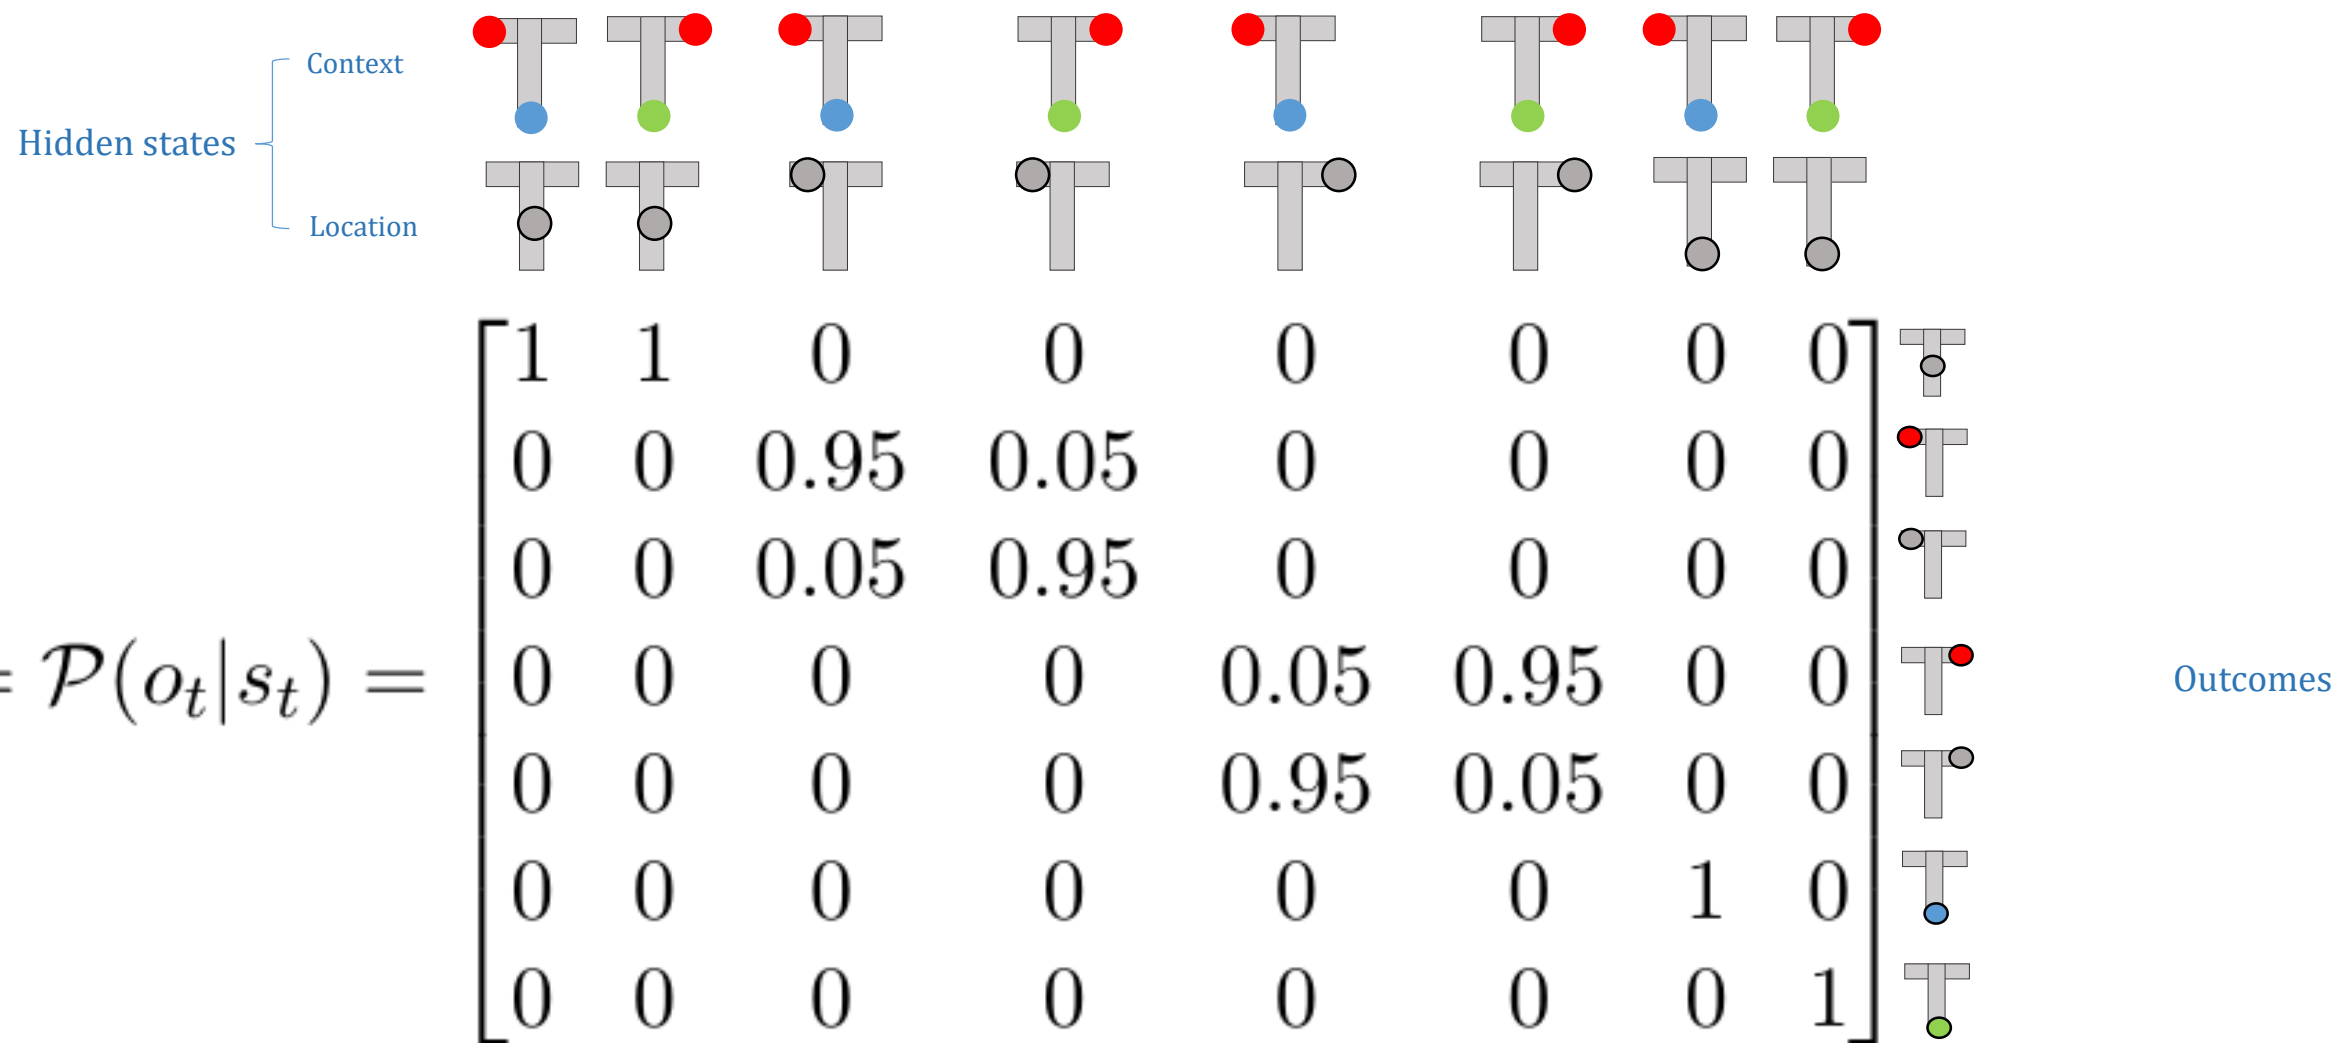
$o \in O$



→ **Toolbox:** ABC of choice behaviour

ABC of choice behaviour:

A – Mapping from hidden states to outcomes



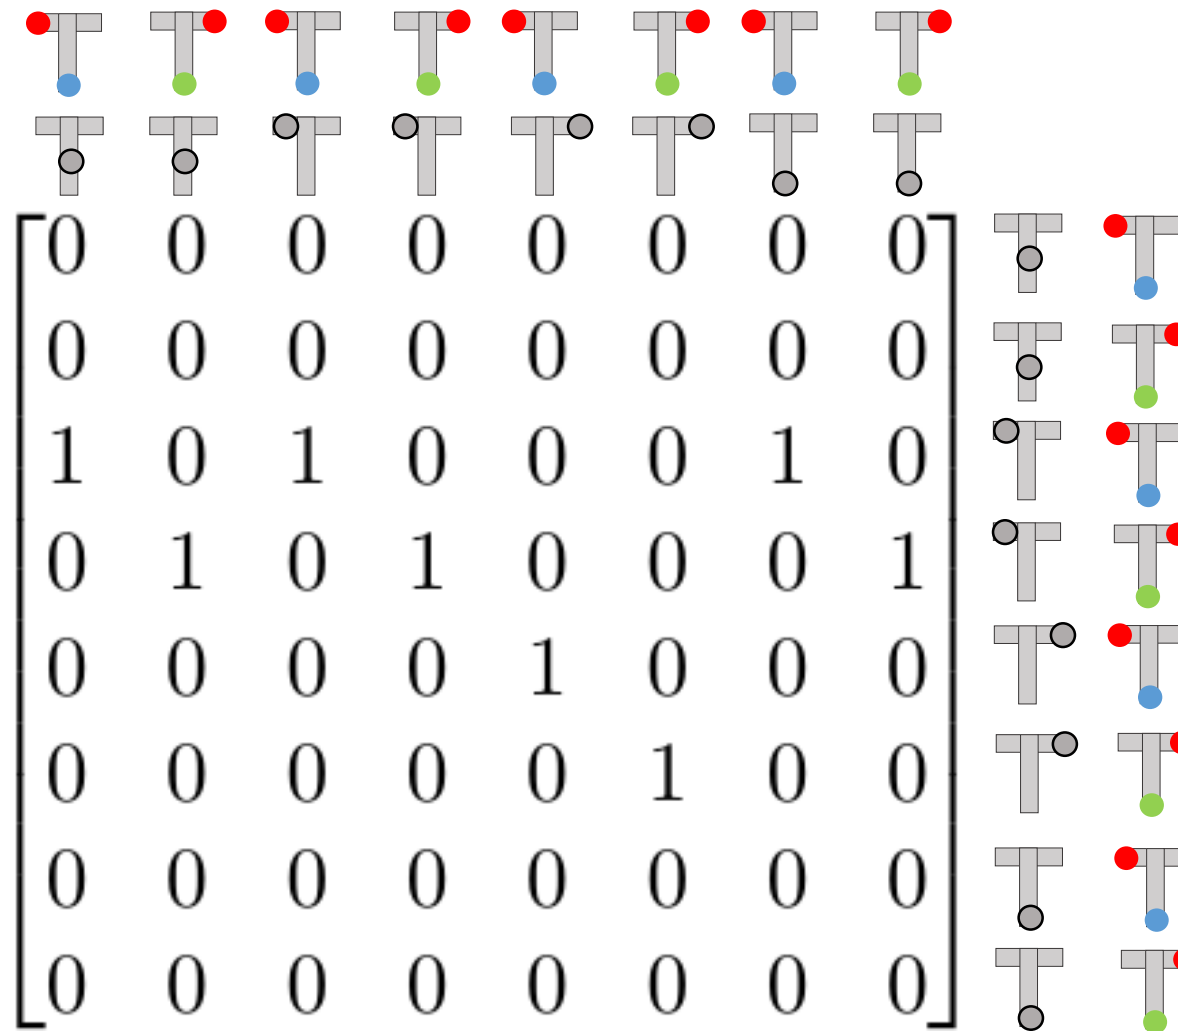
ABC of choice behaviour: B – Transition Probabilities

Hidden states

Location

$$\mathcal{B}\{\text{⌞}\} = \mathcal{P}(s_{t+1}|s_t, \text{⌞}) =$$

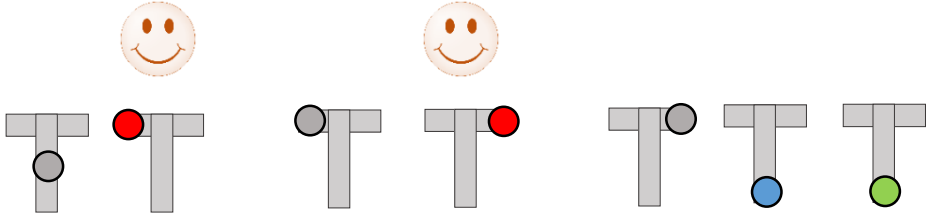
Also: $\mathcal{B}\{\text{TTT}\}$, $\mathcal{B}\{\text{T}\}$, $\mathcal{B}\{\text{TT}\}$...



Hidden states

ABC of choice behaviour:

C – Preferences over Outcomes

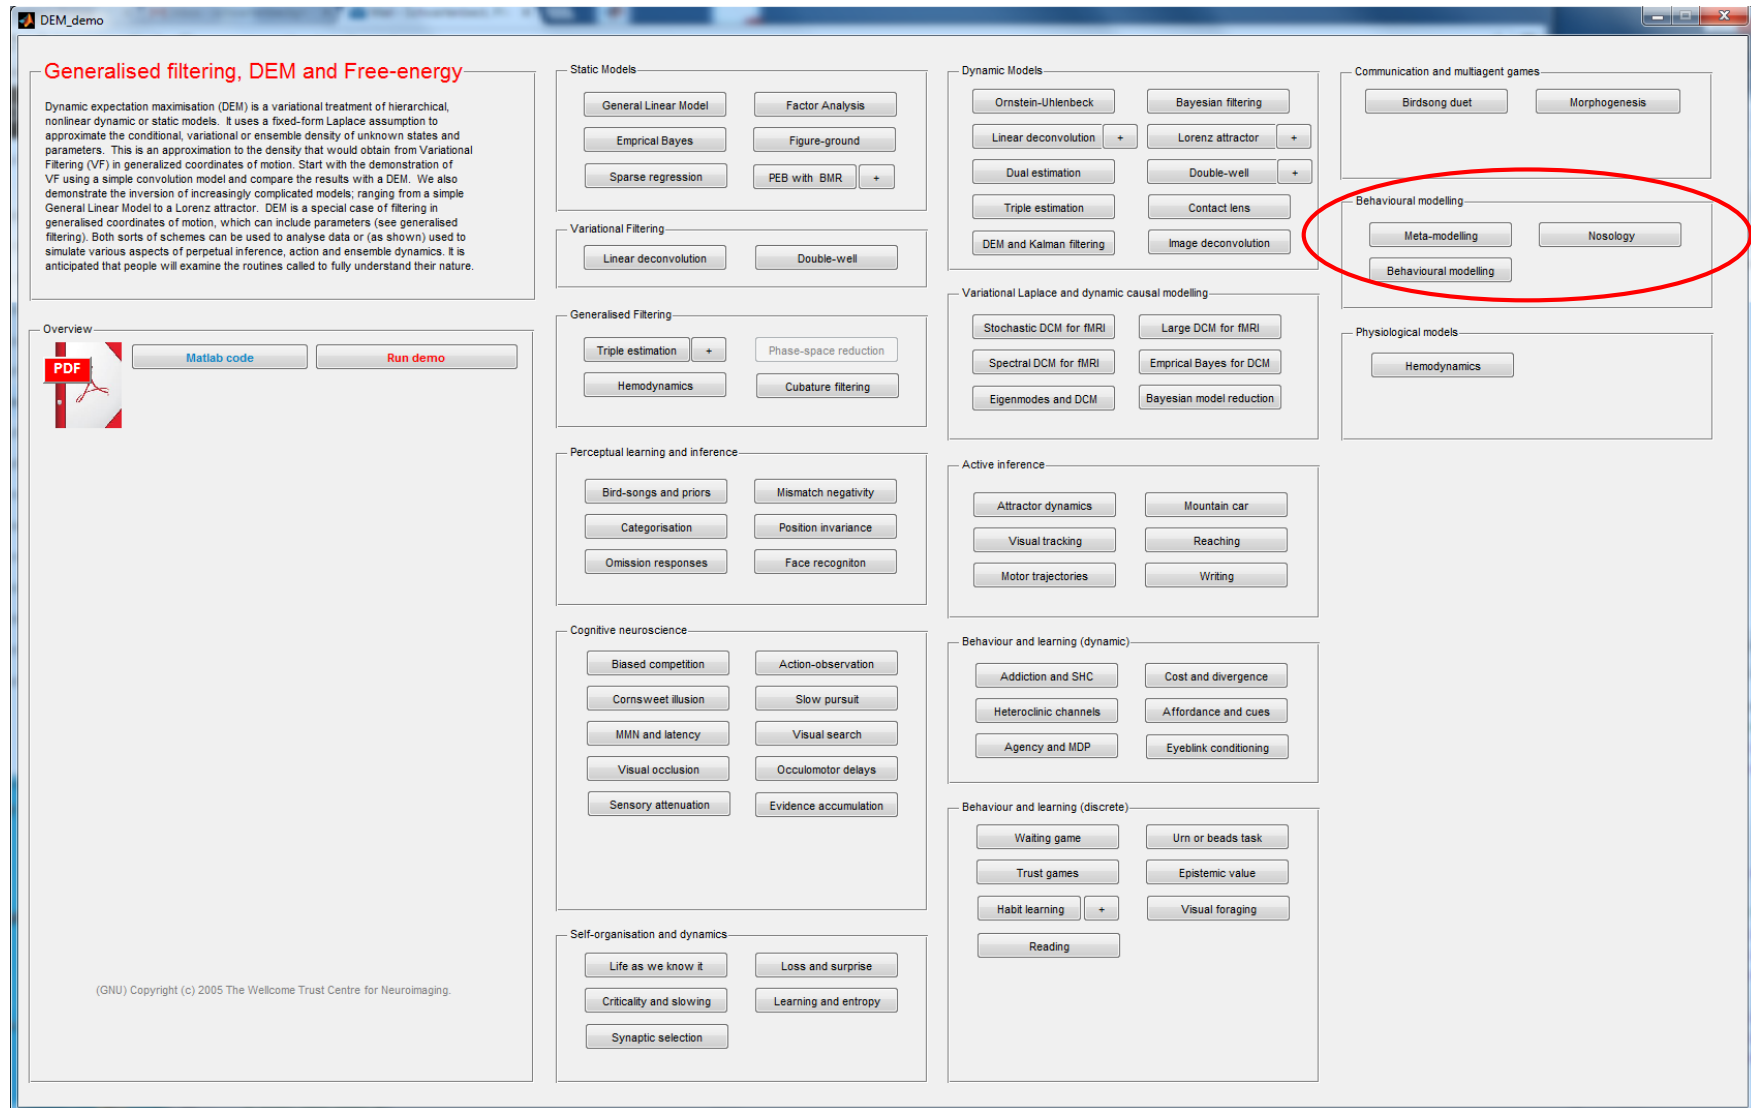


The diagram illustrates a choice task with two options, each represented by a grey T-shaped lever. In the first pair, the left lever has a grey ball and the right lever has a red ball; a smiley face is above the right lever. In the second pair, the left lever has a grey ball and the right lever has a red ball; a smiley face is above the right lever. In the third pair, the left lever has a grey ball and the right lever has a blue ball. In the fourth pair, the left lever has a grey ball and the right lever has a green ball.



$$c = \ln \mathcal{P}(o_t) = [0 \quad 2 \quad -2 \quad 2 \quad -2 \quad 0 \quad 0]^T$$

Outcomes

DEM Toolbox



ABC of choice behaviour: Prior over initial states and allowable policies



- I. ABC of choice behaviour
- II. Simulate data and invert model to obtain subject-specific parameters
- III. Simulate group effect and recover group-specific parameters
- IV. Perform hierarchical Bayesian inference on group effects

```
%% set up and preliminaries: first generate synthetic (single subject) data
=====
rng('default')

% outcome probabilities: A
%-----
% We start by specifying the probabilistic mapping from hidden states
% to outcomes.
%-----
a      = .95;
b      = 1 - a;
A      = [1 1 0 0 0 0 0 0;      % ambiguous starting position (centre)
          0 0 a b 0 0 0 0;      % left arm selected and rewarded
          0 0 b a 0 0 0 0;      % left arm selected and not rewarded
          0 0 0 0 b a 0 0;      % right arm selected and rewarded
          0 0 0 0 a b 0 0;      % right arm selected and not rewarded
          0 0 0 0 0 0 1 0;      % informative cue - reward on right
          0 0 0 0 0 0 0 1];     % informative cue - reward on left

% controlled transitions: B{u}
%-----
% Next, we have to specify the probabilistic transitions of hidden states
% under each action or control state. Here, there are four actions taking the
% agent directly to each of the four locations.
%-----
B{1} = [1 0 0 1; 0 1 0 0; 0 0 1 0; 0 0 0 0]; % move to the middle
B{2} = [0 0 0 0; 1 1 0 1; 0 0 1 0; 0 0 0 0]; % move up left (and check for reward)
B{3} = [0 0 0 0; 0 1 0 0; 1 0 1 1; 0 0 0 0]; % move up right (and check for reward)
B{4} = [0 0 0 0; 0 1 0 0; 0 0 1 0; 1 0 0 1]; % move down (check cue)

for i = 1:4
    B{i} = kron(B{i}, eye(2));
end
```

```
% priors: (utility) C
%-----
% Finally, we have to specify the prior preferences in terms of log
% probabilities. Here, the agent prefers rewarding outcomes
%-----
c = 2;
C = [0 c -c c -c 0 0]';

% now specify prior beliefs about initial state, in terms of counts
%-----
d = kron([1 0 0 0], [1 1])';

% allowable policies (of depth T). These are just sequences of actions
%-----
V = [1 1 1 1 2 3 4 4 4 4
     1 2 3 4 2 3 1 2 3 4];
```

```
%% MDP Structure - this will be used to generate arrays for multiple trials
=====
mdp.V = V; % allowable policies
mdp.A = A; % observation model
mdp.B = B; % transition probabilities
mdp.C = C; % preferred states
mdp.D = d; % prior over initial states
mdp.s = 1; % initial state

mdp.alpha = 2; % precision of action selection
% mdp.alpha = 8; % precision of action selection
mdp.beta = 1; % inverse precision of policy selection

% true parameters
%-----
n = 128; % number of trials
i = rand(1,n) > 1/2; % randomise hidden states over trials
P.beta = log(2);
P.C = log(2);

MDP = mdp;
% MDP.C = mdp.C;
MDP.C = mdp.C*exp(P.C);
MDP.beta = mdp.beta*exp(P.beta);

[MDP(1:n)] = deal(MDP);
[MDP(i).s] = deal(2);
```

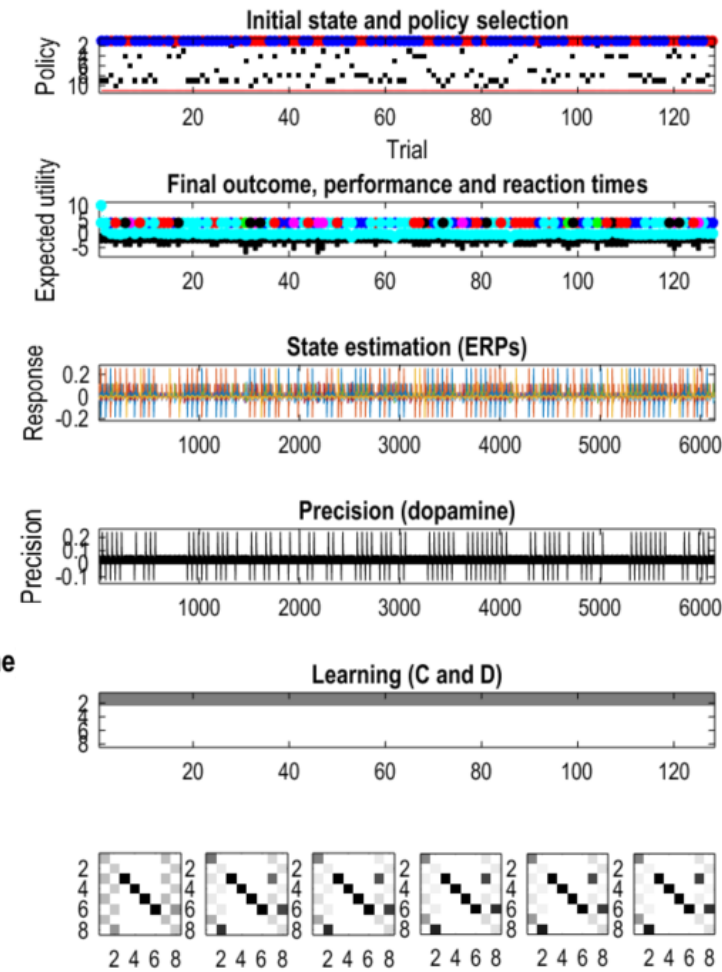
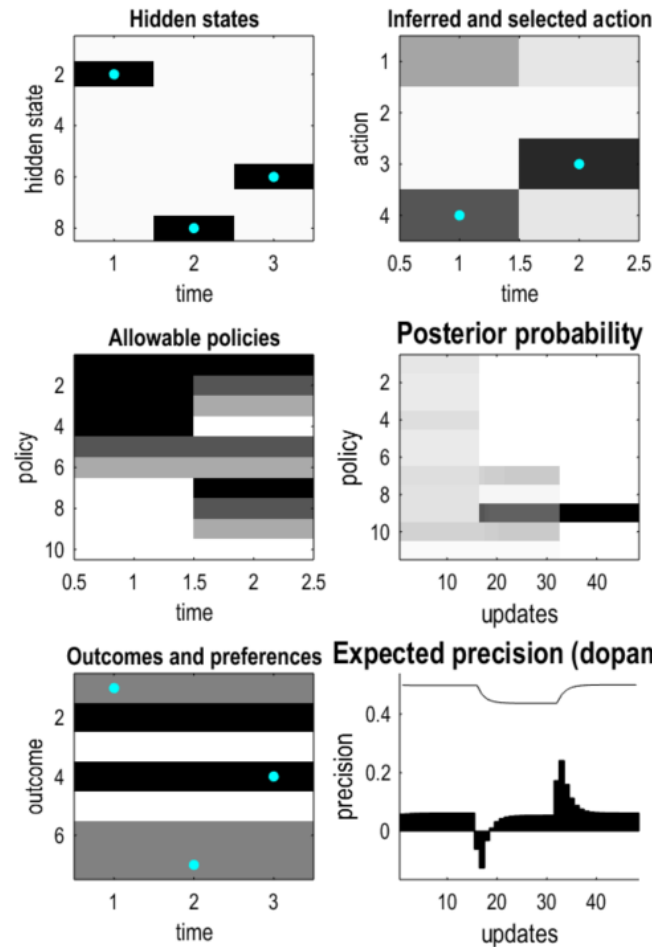
- I. ABC of choice behaviour
- II. **Simulate data** and invert model to obtain subject-specific parameters (**spm_MDP_VB**)
- III. Simulate group effect and recover group-specific parameters
- IV. Perform hierarchical Bayesian inference on group effects

```
%% Solve to generate data
%=====
MDP = spm_MDP_VB(MDP);

% illustrate behavioural responses - single trial
%-----
spm_figure('GetWin','Figure 1a'); clf
spm_MDP_VB_trial(MDP(1));

% illustrate behavioural responses and neuronal correlates over trials
%-----
spm_figure('GetWin','Figure 1b'); clf
spm_MDP_VB_game(MDP);

% This completes the generation of data. We now turn to the estimation of
% subject specific preferences and precision encoded by the parameters
% beta and C. Model parameters here are log scaling parameters that allow
% for increases or decreases in the default prior values.
%-----
```

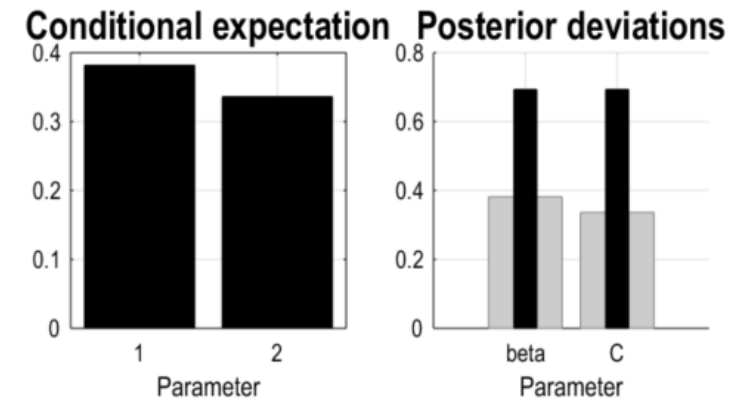
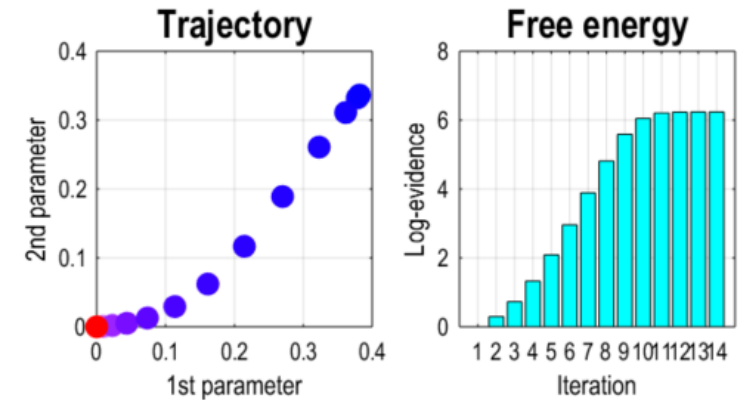


- I. ABC of choice behaviour
- II. Simulate data and **invert model to obtain subject-specific parameters (spm_dcm_mdp)**
- III. Simulate group effect and recover group-specific parameters
- IV. Perform hierarchical Bayesian inference on group effects

```
%% Invert to recover parameters (preferences and precision)
%=====
DCM.MDP = mdp; % MDP model
DCM.field = {'beta', 'C'}; % parameter (field) names to optimise
DCM.U = {MDP.o}; % trial specification (stimuli)
DCM.Y = {MDP.u}; % responses (action)

DCM = spm_dcm_mdp(DCM);

% compare true values with posterior estimates
%-----
subplot(2,2,4), hold on
bar(spm_vec(P), 1/4)
set(gca, 'XTickLabel', DCM.field)
set(gcf, 'Name', 'Figure 2', 'Tag', 'Figure 2')
```



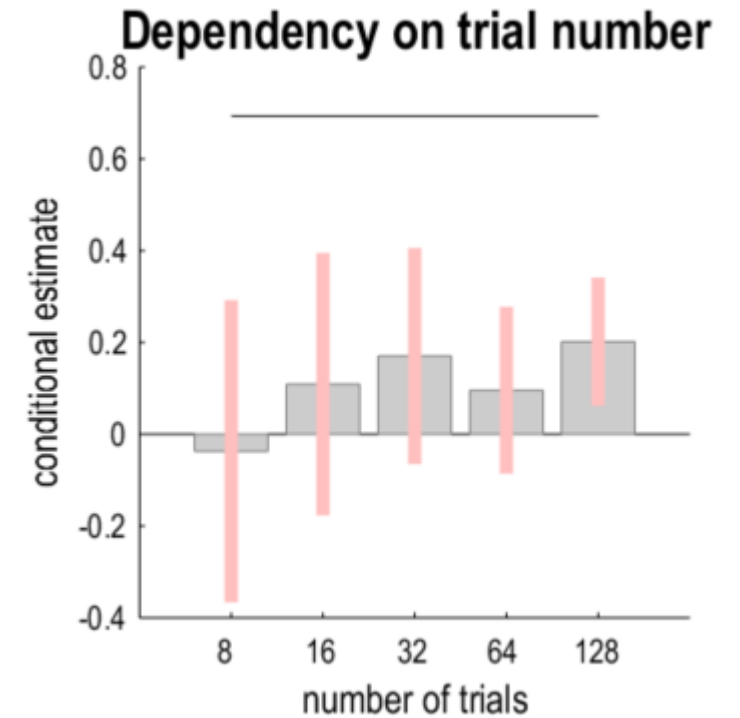
Optimise design parameters (trial number)

```

%% now repeat using subsets of trials to illustrate effects on estimators - design optimisation!
%=====
DCM.field = {'beta'};
n          = [8 16 32 64 128];
for i = 1:length(n)
    DCM.U = {MDP(1:n(i)).o};
    DCM.Y = {MDP(1:n(i)).u};
    DCM   = spm_dcm_mdp(DCM);
    Ep(i,1) = DCM.Ep.beta;
    Cp(i,1) = DCM.Cp;
end

% plus results
%-----
spm_figure('GetWin','Figure 3'); clf
subplot(2,1,1), spm_plot_ci(Ep(:),Cp(:)), hold on
plot(1:length(n),(n - n) + P.beta,'k'), hold off
set(gca,'XTickLabel',n)
xlabel('number of trials','FontSize',12)
ylabel('conditional estimate','FontSize',12)
title('Dependency on trial number','FontSize',16)
axis square

```



- I. ABC of choice behaviour
- II. Simulate data and invert model to obtain subject-specific parameters
- III. Simulate group effect and **recover group-specific parameters (spm_dcm_fit)**
- IV. Perform hierarchical Bayesian inference on group effects

```
% now repeat but over multiple subjects with different beta
%=====

% generate data and a between subject model with two groups of eight
% subjects
%-----
N = 8; % numbers of subjects per group
X = kron([1 1 -1],ones(N,1)); % design matrix
h = 4; % between subject log precision
n = 128; % number of trials
i = rand(1,n) > 1/2; % randomise hidden states

clear MDP
[MDP(1:n)] = deal(mdp);
[MDP(i).s] = deal(2);

reward = zeros(n,size(X,1));

for i = 1:size(X,1)

    % true parameters - with a group difference of one half
    %-----
    beta(i) = X(i,:)*[0; 1/4] + exp(-h/2)*randn; % add random Gaussian effects to group means -> BMR and PEB
    % beta(i) = X(i,:)*[0; 0] + exp(-h/2)*randn; % add random Gaussian effects to group means -> BMR and PEB
    [MDP.beta] = deal(exp(beta(i)));

    % solve to generate data
    %-----
    DDP = spm_MDP_VB(MDP); % realisation for this subject
    DCM.U = {DDP.o}; % trial specification (stimuli)
    DCM.Y = {DDP.u}; % responses (action)
    GCM{i,1} = DCM;

    for kk=1:length(DCM.U)
        if DCM.U{kk}(end)==2 || DCM.U{kk}(end)==4 % outcome 2 or 4 == reward
            reward(kk,i)=1;
        end
    end

    % plot behavioural responses
    %-----
    spm_figure('GetWin','Figure 4'); clf
    spm_MDP_VB_game(DDP);drawnow

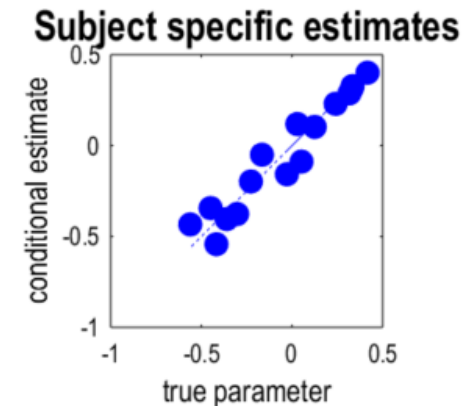
end
```

```
%% Bayesian model inversion
%=====

GCM = spm_dcm_fit(GCM);

% plot subject specific estimates and true values
%-----

spm_figure('GetWin','Figure 4');
subplot(3,1,3)
for i = 1:length(GCM)
    qP(i) = GCM{i}.Ep.beta;
end
plot(beta,beta,':b',beta,qP,':b','MarkerSize',32)
xlabel('true parameter','FontSize',12)
ylabel('conditional estimate','FontSize',12)
title('Subject specific estimates','FontSize',16)
axis square
```

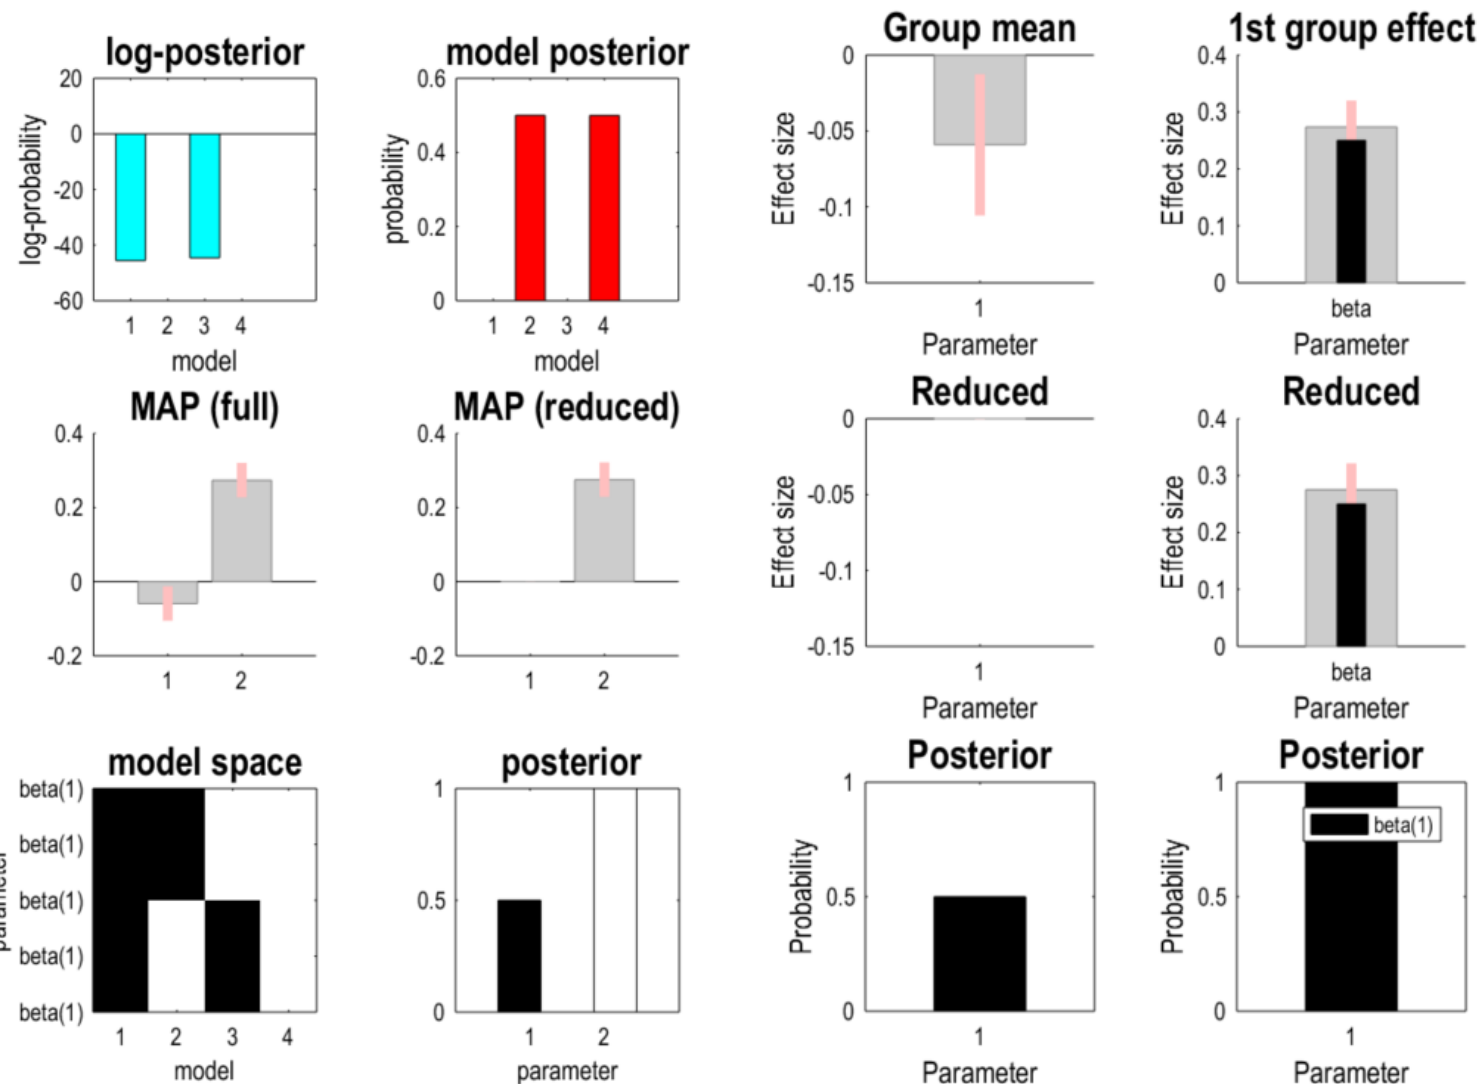


- I. ABC of choice behaviour
- II. Simulate data and invert model to obtain subject-specific parameters
- III. Simulate group effect and recover group-specific parameters
- IV. **Perform hierarchical Bayesian inference on group effects**
(spm_dcm_peb)

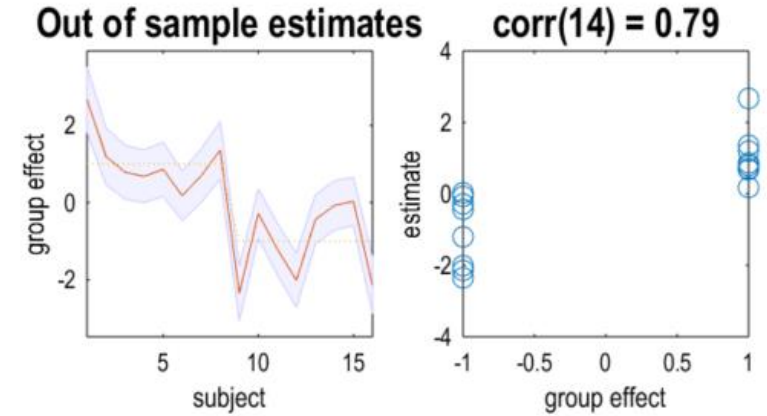
```

%% hierarchical (empirical) Bayes
=====
% second level model
%
M = struct('X',X);
%
% BMA - (second level)
%
PEB = spm_dcm_peb(GCM,M);
BMA = spm_dcm_peb_bmc(PEB);
subplot(3,2,4),hold on, bar(1,1/4,1/4), set(gca,'XTickLabel',DCM.field)
subplot(3,2,2),hold on, bar(1,1/4,1/4), set(gca,'XTickLabel',DCM.field)

```



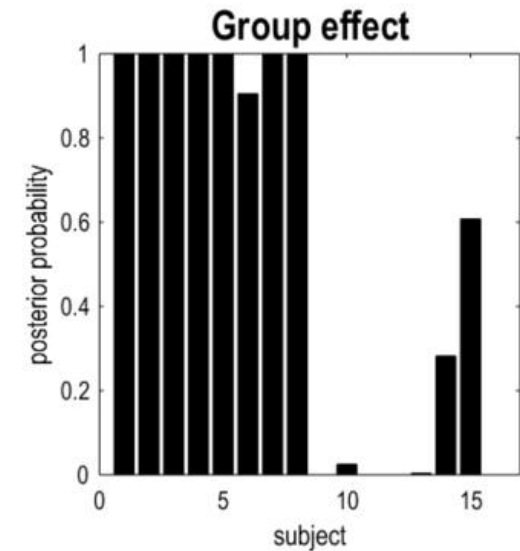
- I. ABC of choice behaviour
- II. Simulate data and invert model to obtain subject-specific parameters
- III. Simulate group effect and recover group-specific parameters
- IV. Perform hierarchical Bayesian inference on group effects and **leave-one-out cross-validation**
(`spm_dcm_loo`)



```
%% posterior predictive density and cross validation
```

```
%=====
```

```
spm_dcm_loo(GCM,M,DCM.field);
```



Active Inference and MDP toolbox - Summary

Active inference introduces concept of surprise minimisation to modelling (choice) behaviour

- Central role of *subjective* and *objective* generative models

Using the toolbox essentially requires defining the state space of a task

- ABC of choice behaviour

Computational phenotyping in Psychiatry:

- Routines for simulating data, inverting models and inferring group effects freely available in SPM (DEM toolbox)
- Test empirical predictions, such as role of precision, expected utility vs. surprise minimisation or Bayesian model reduction

Active Inference and MDP toolbox - Summary

1. Why is this model useful?
 - (Choice) behaviour as probabilistic inference
 - Distinction between inference and learning
2. Where can you use it?
 - Any Inference or planning problem, particularly choice behaviour
 - Any learning problem (more recent version)
3. Where can't you use it?
 - Continuous problems (cf., predictive coding)
4. What do we like about it?
 - Applies normative theory of belief-updating (*Free Energy Principle*) to choice behaviour
 - Highlights importance of prior beliefs to understand (suboptimal) behaviour
 - Highlights necessity to build and update generative models of the world
5. What are the most common mistakes?
 - Definition of state space

Active Inference and MDP - Overview

Theory

- ‘Anatomy of choice’: MDP as surprise minimisation (*Friston et al., 2013*)
- Epistemic value: exploration versus exploitation (*Friston et al., 2015*)
- Active inference and learning (*Friston et al., 2016*)
- Artificial insight (*Friston et al., in prep*)
- + Worked example for computational phenotyping (*Schwartenbeck & Friston, 2016*) and hierarchical Bayesian inference and Bayesian model reduction (*Friston et al., 2015*)

Experimental work

- Role of dopamine, model-based fMRI (*Cerebral Cortex, 2014*)
- Interpersonal inference & trust games (*Front Comp Neurosci, 2014*)
- Evidence accumulation (*Neural Computations, 2015*)
- Economic decision theory (*Scientific Reports, 2015*)
- Scene construction & visual foraging (*Front Comp Neurosci, 2016*)
- Model averaging and structure learning, model-based fMRI (in prep)
- ...

Thanks to

Karl Friston (UCL)

Thomas FitzGerald (UEA)

Christoph Mathys (UCL)

Martin Kronbichler (CCNS)

Ray Dolan (UCL)

Thank you for listening!